

why added 1023 in exponnet, explain why they didnt go for signed implementation and decided to gor unsigned by shifting 2 power 11 - 1 i.e 1023 places

The addition of **1023 to the exponent in IEEE 754 double-precision (64-bit) floating-point representation** is done because of a technique called "**exponent biasing**."

Why Add 1023 to the Exponent? (Biasing Explained)

- In IEEE 754 double precision, the exponent field is **11 bits** (so it can store values 0–2047).
- **Floating-point numbers can have both positive and negative exponents**, allowing them to represent both very large and very small numbers (e.g., 2^{+300} and 2^{-300}).
- However, if you simply stored the exponent as a *signed* value using two's complement, zero (the midpoint) would not be symmetrical or convenient in binary comparison and sorting operations.

So, instead, they use a "bias":

- The "true" exponent value (after normalizing the binary number) is shifted **upwards by a fixed amount** before being stored.
- For doubles (11 exponent bits), the bias is $2^{11-1} - 1 = 1023$. That is, the stored exponent is:
$$\text{stored exponent} = \text{actual exponent} + 1023$$
- For example, to represent an actual exponent of 0 (i.e., scaling by 2^0), the bit pattern stored in the exponent field is $1023 = 0111111111_2$.

Why Not Use Signed Exponent Representation?

- **Binary storage and comparison:** Using only unsigned binary values for the exponent makes it much easier for computers to compare floating-point numbers with ordinary unsigned integer logic. This helps with hardware implementation, sorting, and range checks.
- **All zero bits and all one bits are reserved for special cases:**
 - Exponent = 0 (all zeros) is reserved for *denormal/subnormal* numbers and zeros.
 - Exponent = 2047 (all ones) is used for *infinities* and *NaN* (Not a Number).

If you used two's complement or another signed scheme, you'd need more complex hardware logic to handle those special cases and to decode and compare exponent fields.

- **Uniformity and simplicity:** A bias ensures that all exponents, including negative ones, can be represented as simple, non-negative binary numbers, which are easier to implement in

hardware.

How is the Bias Calculated?

- For an exponent field of n bits, the bias is $2^{n-1}-1$.
 - For double precision (11 bits): $2^{10}-1 = 1023$
 - For single precision (8 bits): $2^7-1 = 127$

In summary:

- **1023 is added to the exponent so all values stored are non-negative, making hardware simpler and special cases easier to represent and detect.**
- **Signed representation was not used** to keep things efficient, compatible with binary comparison, and to avoid overlap with special reserved values. The biasing scheme (shifting the exponent by 1023 for 64-bit floats) is a deliberate design to maximize precision, minimize complexity, and cover the needed dynamic range.